

POSIX.4 Message Queue

- A POSIX.4 message queue is a **priority queue** of discrete messages
- It allows two or more processes to communicate and provides messages' prioritization.
- Within a process, shared variables with mutex is an efficient way to communicate
- What about communication among different processes? Different processes can communicate with message queues. Their use is easier than shared memory regions and avoids to use mutex across processes' boundary.

Creating/opening a message queue

```
#include <mqueue.h>
#define Q_NAME ``/CS431queue'' // the message queue name should start with '/'
struct mq_attr  mq_attr;
mqd_t  mymq;
//set up the message queue attributes
memset(&mq_attr, 0, sizeof(mq_attr)); //clear up any default settings
mq_attr.mq_maxmsg = 10;    // maximum number of messages stored in the message queue
mq_attr.mq_msgsize = 128;  // size of messages expressed as number of bytes
```

Creating/opening a message queue

read/write permissions ← `mq_open(Q_NAME, O_CREAT | O_RDWR, S_IRWXU, &mq_attr);` → user permissions (i.e. unix owner, group, others)

↙
The system will create the message queue if it does not already exist.

Be aware that the system will not return any error if the queue existed already!

If **O_CREAT** | **O_EXCL** are used, the system will return an error if the queue already exists

- The queue_name must be constructed like a normal file pathname, **BUT:**
- To be portable across different platforms, the queue_name should start with `/'
- To be portable across different platforms, the queue_name should contain no other `/' characters

Creating/opening a message queue

```
#include <mqueue.h>
```

```
//create a message queue object
```

```
mymq = mq_open(Q_NAME, O_CREAT | O_RDWR, S_IRWXU, &mq_attr);
```

- oflag → **O_CREAT** causes the message queue object to be created if necessary (see man pages for additional details)
- oflag → you can set flags for read/write permissions (see man pages for additional details)
- mode → it allows to set user permissions (i.e. unix owner, group, others) (see man pages for additional details)
- Remember: **mq_open** operates on a name that may not exist in the file system. So, message queues may not show up in the output of a command like *ls*.

Example with Message Queues

//program1: user interface, create a message queue and send waveform spec

```
void main(void)
```

```
{...
```

```
long buffer[4];
```

```
memset(buffer, 0, sizeof(buffer)); //clean up the buffer space
```

```
//buffer[0] for msg seq #, buffer[1] for min_volt, buffer[2] for max_volt, buffer[3] for freq
```

```
struct mq_attr  wave_q_attr
```

```
//set up the message queue attributes
```

```
memset(&wave_q_attr, 0, sizeof(wave_q_attr)); //clear up any default settings
```

```
wave_q_attr.mq_maxmsg = 10;
```

```
wave_q_attr.mq_msgsize = sizeof(buffer);
```

```
...
```

```
(to be continued)
```

It tells the system not to suspend when you try to receive on an empty queue or send to a queue where there is no space!

Example: Sending Messages

(continue)

//WAVE_SPEC_Q is our name for the message queue. It is GLOBAL across all processes.

mqd_t wave_q_id;

wave_q_id = mq_open("/WAVE_SPEC_Q", O_WRONLY | O_CREAT | O_NONBLOCK, 0600, &wave_q_attr);

// 600 = queue permission, same as file permission: owner(rw-), group(- - -), others(- - -)

while(command != 'q') {

//ask user if (s)he wants to give new spec or quit

// if 'q' is received, set the message sequence number to -1, otherwise, increment it by 1

//get new waveform spec data from user and put the waveform spec into the buffer

buffer_length = sizeof(buffer);

status = mq_send(wave_q_id, (char *) buffer, buffer_length, prio);

//prio = the priority assigned to msg, a rule of thumb is to use the sender thread priority.

//Real-time messages are prioritized. Not FIFO. Same priority msg are kept in FIFO order

}

//clean up

status = mq_close(wave_q_id); //like closing a file, the message queue STILL lives in the system

status = mq_unlink("/WAVE_SPEC_Q"); //destroy the message queue after everyone closes the queue

}

Example: Receiving Messages

//program 2: create a receiving queue for waveform spec msg, and display the wave form

```
void main(void)
```

```
{    long buffer[4];
```

```
    struct mq_attr  wave_q_attr
```

```
    //set up the same message queue attribute
```

```
    memset(&wave_q_attr, 0, sizeof(wave_q_attr));
```

```
    wave_q_attr.mq_maxmsg = 10;
```

```
    wave_q_attr.mq_msgsize = sizeof(buffer);
```

```
    ...
```

```
    //set up the buffer to store the waveform spec,
```

```
    //buffer[0] = msg seq #, buffer[1] = min_volt, buffer[2] = max_volt, buffer[3] = freq
```

```
    memset(buffer, 0, sizeof(buffer)); //clean up the buffer space
```

```
    mqd_t wave_q_id;
```

```
    wave_q_id = mq_open("/WAVE_SPEC_Q", O_RDONLY | O_CREAT | O_NONBLOCK,  
                        0600, &wave_q_attr);
```

```
    while (buffer[0] != -1) {    // execute the loop if the message seq # != -1
```

```
        buffer_length = sizeof(buffer);
```

```
        status = mq_receive(wave_q_id, (char *) buffer, buffer_length, &prior);
```

```
        //& prior: priority of the message that you just read.
```

```
        //check message sequence number, get the wave form spec from buffer
```

```
        //use posix timer to periodically send out the voltages to display the waveform
```

```
    } }
```

It informs the system
how large is the
receiver buffer!

State of Message Queues

- Life span of a message queue
 - Once a message queue is created, like a file, its life is longer than the life of the program that creates it, unless it is unlinked.
 - This feature permits debugging after a program fails and allows for different programs to read and write to shared message queues.
 - Unlike a file, a message queue will disappear if the system reboots.
 - Message queues are inherited when a process **forks**. In addition, when a process calls **exit** or **exec**, message queues are closed implicitly.
- Quiz: in our previous example, what would happen if we use `cntr_c`, instead of using the “q”, to terminate the program, and then restart the program again?

POSIX Example: mq_master.c

```
#include <signal.h>
#include <time.h>
#include <fcntl.h> /* For O_* constants */
#include <sys/stat.h> /* For mode constants */
#include <mqueue.h>
#include <sched.h>
#include <stdio.h>
#include <stdlib.h>

#define MY_MQ "/my_mq"

sigset_t signal_set;

void init_timer(long period)
{
    timer_t timer;
    struct itimerspec timer_spec;

    // Create a timer.
    timer_create(CLOCK_REALTIME, NULL, &timer);

    // Specify the timer's period.
    timer_spec.it_value.tv_sec = 0;
    timer_spec.it_value.tv_nsec = 1; /* Start immediately */
    timer_spec.it_interval.tv_sec = 0;
    timer_spec.it_interval.tv_nsec = period;

    // Setup a signal set for sigwait() to wait for SIGALRM
    sigemptyset(&signal_set);
    sigaddset(&signal_set, SIGALRM);
    sigprocmask(SIG_BLOCK, &signal_set, NULL);

    // Setup the timer's period.
    timer_settime(timer, 0, &timer_spec, NULL);
}
```

POSIX Example: mq_master.c

```
/* ***** MASTER TASK *****/
```

```
int main(void)
{
    mqd_t queue;
    struct mq_attr queue_attr;

    // Open the message queue read-only and non-
    blocking.
    queue_attr.mq_maxmsg = 10;
    queue_attr.mq_msgsize = sizeof(unsigned long);
    queue = mq_open(MY_MQ, O_RDONLY |
        O_CREAT | O_NONBLOCK, 0600, &queue_attr);
    // Setup scheduler
    struct sched_param spar;
    spar.sched_priority =
        sched_get_priority_max(SCHED_FIFO);
    sched_setscheduler(0, SCHED_FIFO, &spar);
    // Setup timer
    init_timer(20 * 1000000);
```

```
    while (1)
    {
        unsigned long mesg;
        sigwaitinfo(&signal_set, NULL);
        printf("--- New period ---\n");
        while (mq_receive(queue, (char *)&mesg,
            sizeof(mesg), NULL) >= 0)
        {
            printf("Mesg: [%ld]\n", mesg);
        }
        printf("-----\n");
    }

    return EXIT_SUCCESS;
}
```

POSIX Example: mq_slave.c

```
#include <signal.h>
#include <time.h>
#include <fcntl.h> /* For O_* constants */
#include <sys/stat.h> /* For mode constants */
#include <mqueue.h>
#include <sched.h>
#include <stdio.h>
#include <errno.h>
#include <stdlib.h>
```

```
#define MY_MQ "/my_mq"
#define SLAVE_PRIORITY 1
#define SLAVE_PERIOD 10
```

```
sigset_t signal_set;
```

```
void init_timer(long period)
{
    timer_t timer;
    struct itimerspec timer_spec;

    // Create a timer.
    timer_create(CLOCK_REALTIME, NULL, &timer);

    // Specify the timer's period.
    timer_spec.it_value.tv_sec = 0;
    timer_spec.it_value.tv_nsec = 1; /* Start immediately */
    timer_spec.it_interval.tv_sec = 0;
    timer_spec.it_interval.tv_nsec = period;

    // Setup a signal set for sigwait() to wait for SIGALRM
    sigemptyset(&signal_set);
    sigaddset(&signal_set, SIGALRM);
    sigprocmask(SIG_BLOCK, &signal_set, NULL);

    // Setup the timer's period.
    timer_settime(timer, 0, &timer_spec, NULL);
}
```

POSIX Example: mq_slave.c

```
/* ***** SLAVE TASK ***** */
int main(void)
{
    mqd_t queue;
    queue = mq_open(MY_MQ, O_WRONLY);
    // Setup scheduler
    struct sched_param spar;
    spar.sched_priority =
        sched_get_priority_min(SCHED_FIFO) +
        SLAVE_PRIORITY;
    if(sched_setscheduler(0, SCHED_FIFO, &spar) < 0)
    {
        printf("Unable to set SCHED_FIFO. %d\n", errno);
        exit(EXIT_FAILURE);
    }
    // Setup timer
    init_timer(SLAVE_PERIOD * 1000000);
    unsigned long mesg = SLAVE_PRIORITY;
```

```
while (1)
{
    sigwaitinfo(&signal_set, NULL);
    if (mq_send(queue, (char *)&mesg,
        sizeof(mesg), SLAVE_PRIORITY) < 0)
    {
        printf("Unable to send. %d\n", errno);
    }
}

return EXIT_SUCCESS;
}
/**** END - SLAVE TASK *** */
```

POSIX Example: mq_slave2.c

```
#include <signal.h>
#include <time.h>
#include <fcntl.h> /* For O_* constants */
#include <sys/stat.h> /* For mode constants */
#include <mqueue.h>
#include <sched.h>
#include <stdio.h>
#include <errno.h>
#include <stdlib.h>
```

```
#define MY_MQ "/my_mq"
#define SLAVE_PRIORITY 2
#define SLAVE_PERIOD 10
```

```
sigset_t signal_set;
```

```
void init_timer(long period)
{
    timer_t timer;
    struct itimerspec timer_spec;

    // Create a timer.
    timer_create(CLOCK_REALTIME, NULL, &timer);

    // Specify the timer's period.
    timer_spec.it_value.tv_sec = 0;
    timer_spec.it_value.tv_nsec = 1; /* Start immediately */
    timer_spec.it_interval.tv_sec = 0;
    timer_spec.it_interval.tv_nsec = period;

    // Setup a signal set for sigwait() to wait for SIGALRM
    sigemptyset(&signal_set);
    sigaddset(&signal_set, SIGALRM);
    sigprocmask(SIG_BLOCK, &signal_set, NULL);

    // Setup the timer's period.
    timer_settime(timer, 0, &timer_spec, NULL);
}
```

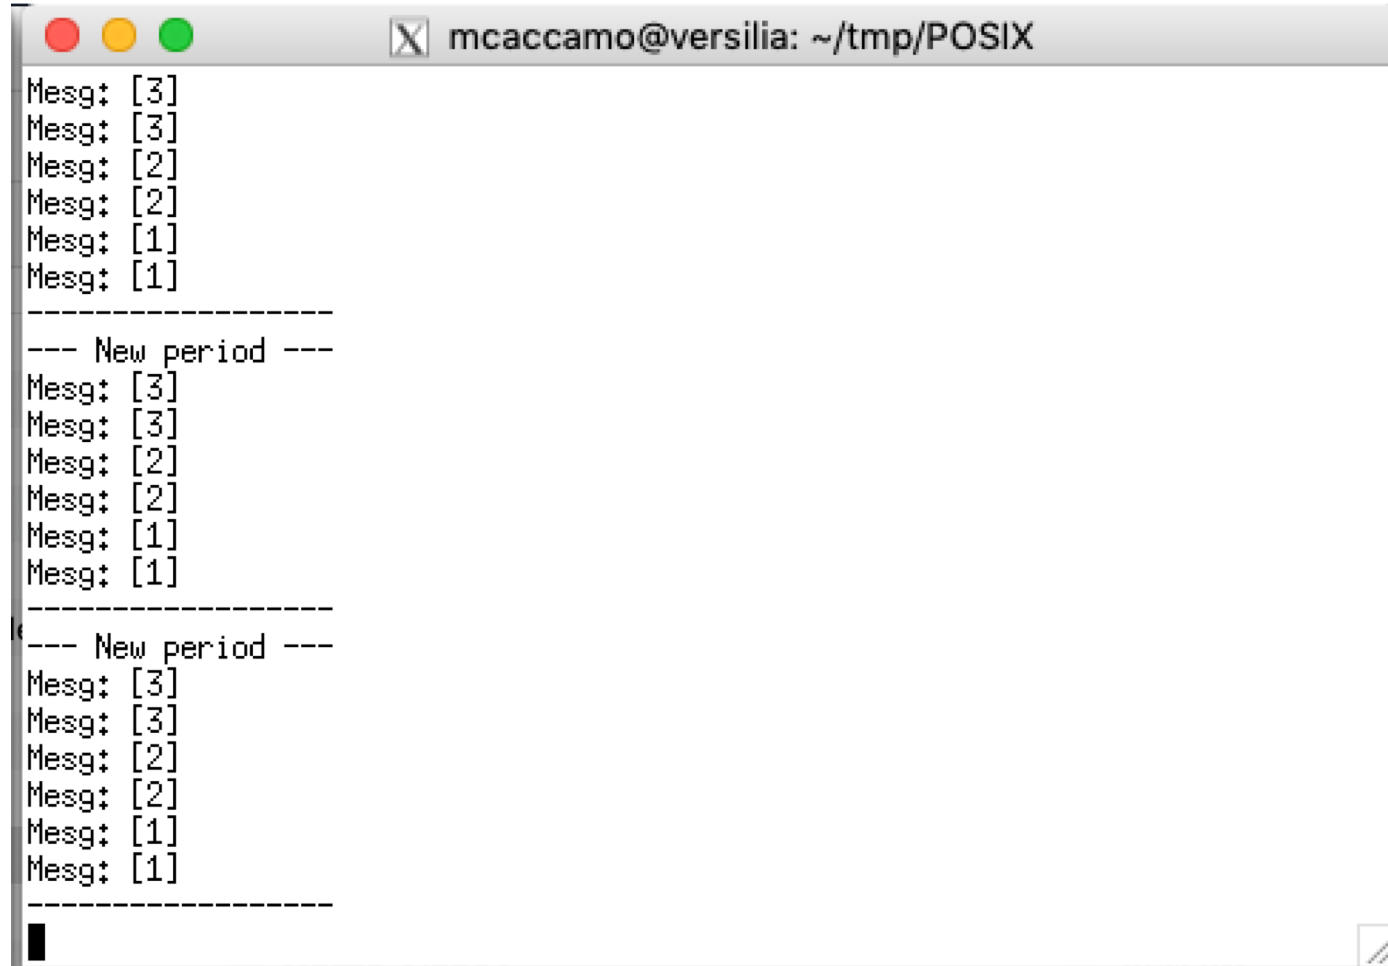
POSIX Example: mq_slave2.c

```
/* ***** SLAVE TASK *****/
int main(void)
{
    mqd_t queue;
    queue = mq_open(MY_MQ, O_WRONLY);
    // Setup scheduler
    struct sched_param spar;
    spar.sched_priority =
        sched_get_priority_min(SCHED_FIFO) +
        SLAVE_PRIORITY;
    sched_setscheduler(0, SCHED_FIFO, &spar);
    // Setup timer
    init_timer(SLAVE_PERIOD * 1000000);
    unsigned long mesg = SLAVE_PRIORITY;
```

```
    while (1)
    {
        sigwaitinfo(&signal_set, NULL);
        if (mq_send(queue, (char *)&mesg,
            sizeof(mesg), SLAVE_PRIORITY) < 0)
        {
            printf("Unable to send. %d\n", errno);
        }
    }

    return EXIT_SUCCESS;
}
/**** END - SLAVE TASK *** */
```

POSIX Example: run



A terminal window titled "mcaccamo@versilia: ~/tmp/POSIX" displays the output of a program. The output consists of three identical blocks, each separated by a dashed line. Each block starts with "---- New period ----" followed by six lines of "Mesg: [3]", "Mesg: [3]", "Mesg: [2]", "Mesg: [2]", "Mesg: [1]", and "Mesg: [1]". The terminal window has a standard macOS-style title bar with red, yellow, and green buttons on the left. A cursor is visible at the bottom left of the terminal area.

```
mcaccamo@versilia: ~/tmp/POSIX
Mesg: [3]
Mesg: [3]
Mesg: [2]
Mesg: [2]
Mesg: [1]
Mesg: [1]
-----
---- New period ----
Mesg: [3]
Mesg: [3]
Mesg: [2]
Mesg: [2]
Mesg: [1]
Mesg: [1]
-----
---- New period ----
Mesg: [3]
Mesg: [3]
Mesg: [2]
Mesg: [2]
Mesg: [1]
Mesg: [1]
-----
```